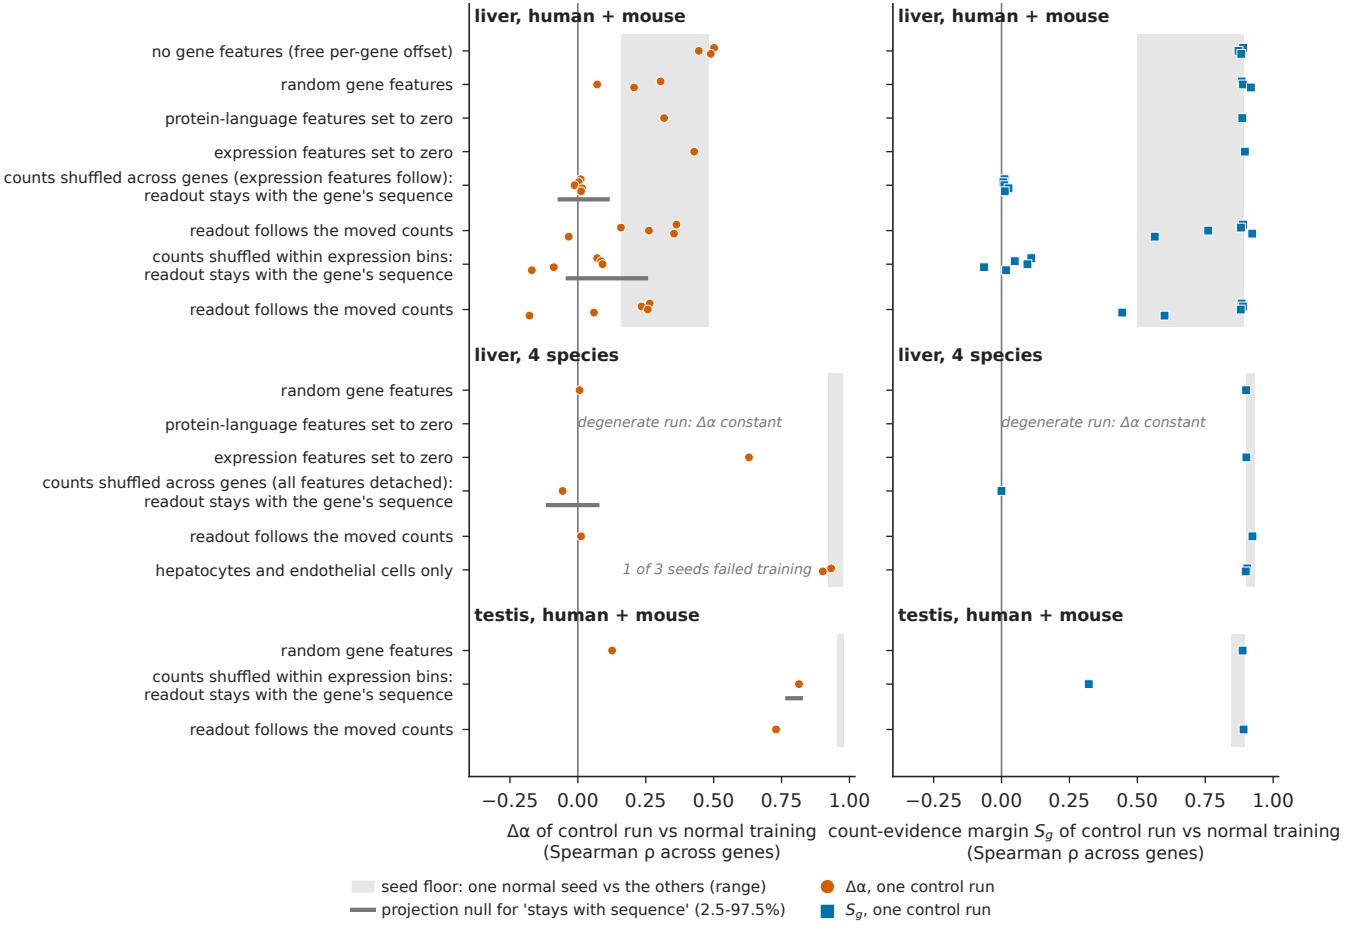
liver, human + mouse
no gene features (free per-gene offset)
random gene features
protein-language features set to zero
expression features set to zero
counts shuffled across genes (expression features follow): readout stays with the gene's sequence
readout follows the moved counts
counts shuffled within expression bins: readout stays with the gene's sequence
readout follows the moved counts
liver, 4 species
random gene features
protein-language features set to zero
degenerate run: Δα constant
expression features set to zero
counts shuffled across genes (all features detached): readout stays with the gene's sequence
readout follows the moved counts
hepatocytes and endothelial cells only
1 of 3 seeds failed training
testis, human + mouse
random gene features
counts shuffled within expression bins: readout stays with the gene's sequence
readout follows the moved counts
−0.25 0.00 0.25 0.50 0.75 1.00
Δα of control run vs normal training (Spearman ρ across genes)
count-evidence margin S_g of control run vs normal training (Spearman ρ across genes)
seed floor: one normal seed vs the others (range)
projection null for 'stays with sequence' (2.5-97.5%)
Δα, one control run
S_g, one control run